

Computer Vision



01

Computer Vision

Introduction
Evaluation of models

02

Transfer Learning & Data Augmentation

State of the art models

03

Robustness

Adversarial Attacks

04

Generative Models

GAN & VEE



Olivier Randavel

Data Scientist

@ olivier.randavel@gmail.com



Hippolyte Mayard

Data Scientist

@ hippolyte.mayard@dauphine.eu



Team

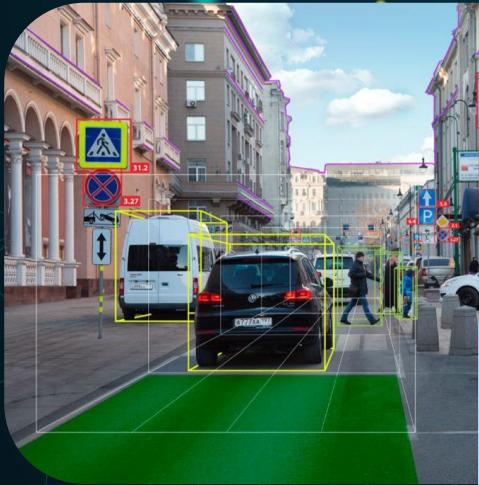
Computer Vision

01

Introduction
Evaluation of models

Applications

Autonomous Vehicles



Face Recognition



Input
Chest X-Ray Image



Radiology

Aims

Race
identification

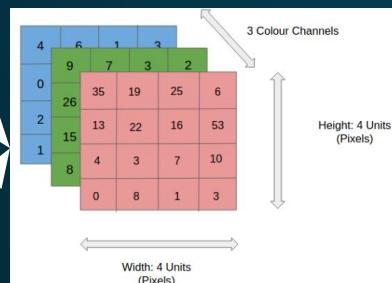
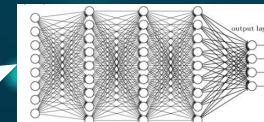


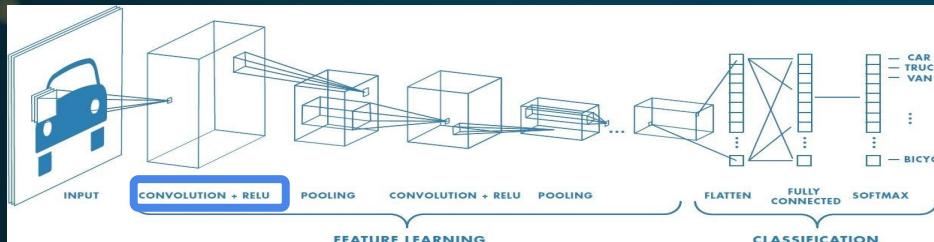
Image to numbers using pixel representations

Deep Neural
network

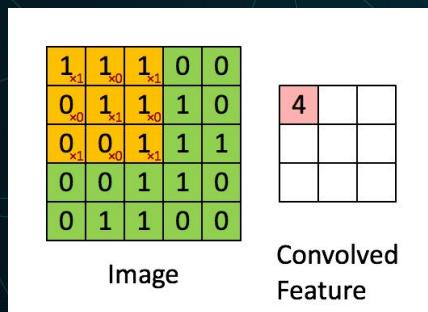


Classification

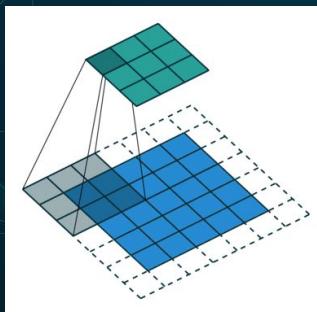
Neural Network



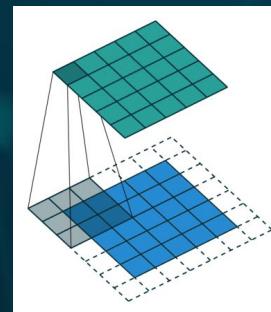
The aim is to train specialised weights in order to detect shape, colors through convolutional layers.



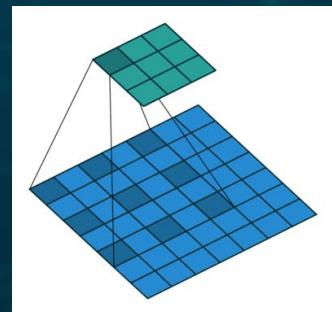
Kernel : reduce dimension



Stride : detect edges

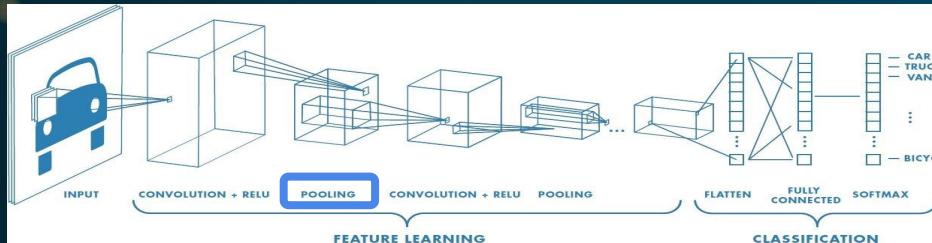


Padding: keep information from borders. Increase or keep same length

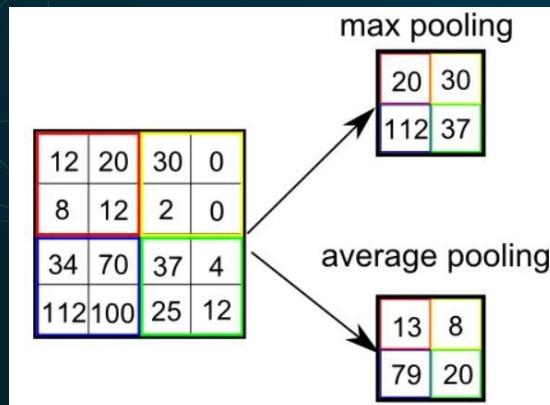


Dilation: segmentation

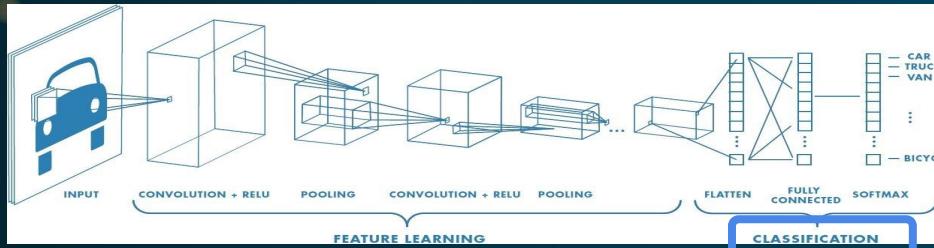
Neural Network



The pooling layer aims to reduce the size of the picture. This helps to reduce the learning time.



Neural Network



The learning process is done then the classification process starts. This is a basic neural network :

- The flatten layer transforms a 3D space into 1D. The picture is a flatten vector
- The fully connected layer uses to extract learnings from all previous layers
- The softmax function to classify pictures

○ SOFTMAX PROBABILITIES

$$S(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

→ $p = 0.7$
→ $p = 0.2$
→ $p = 0.1$

Model Evaluation

Minimise the loss



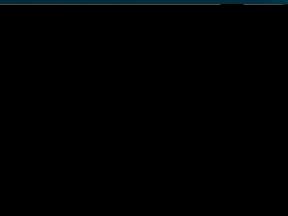
Confusion Matrix

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) <i>Type II Error</i>	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) <i>Type I Error</i>	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
	Precision	Positive Predictive Value $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

Improve the accuracy



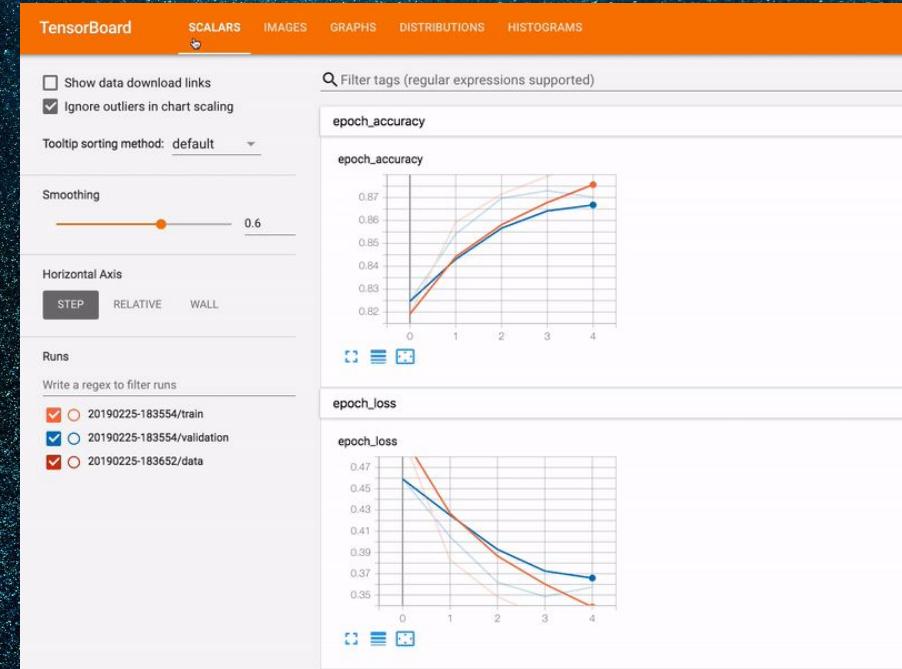
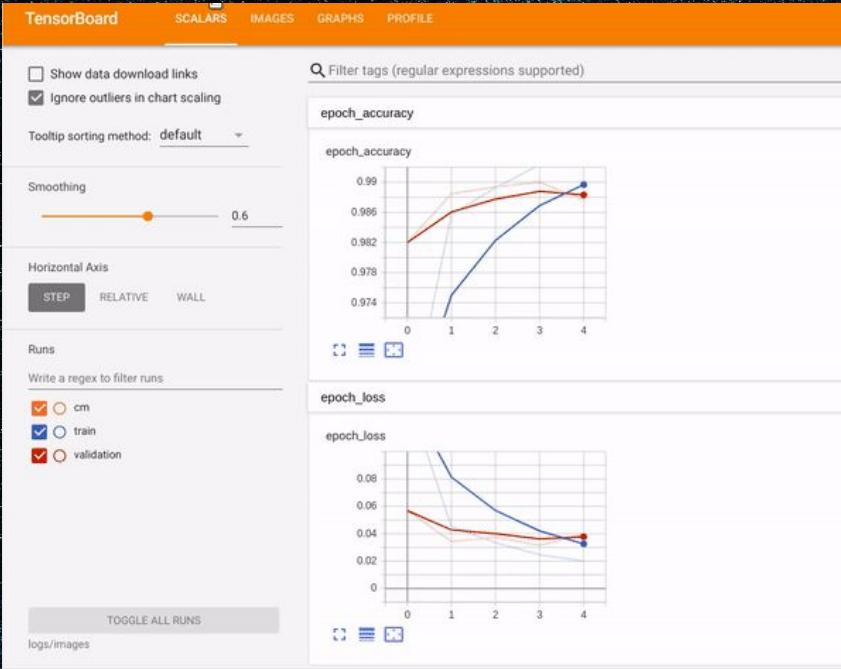
Learn from weights



Try it yourself



Use Tensorboard



Transfer Learning & Data Augmentation

02

State of the art models

Why do we use transfer learning ?

Learning on large amount
of data can cost a lot



Learning on large amount
of data takes times



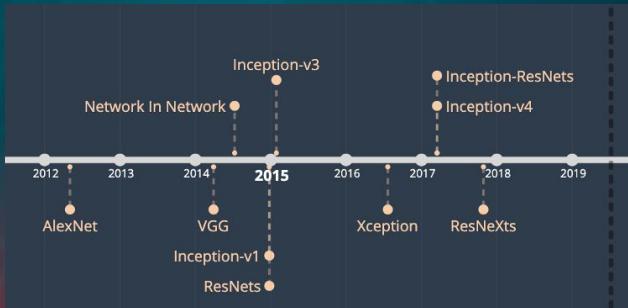
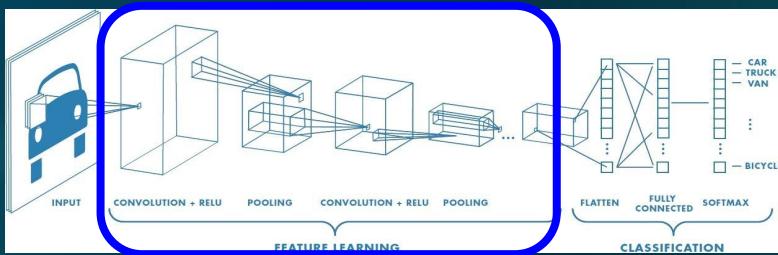
It is preferable when the
model need to be trained
regularly



“Situation where what has been learned in one setting is exploited to improve generalization in another setting.”

Deep Learning, Goodfellow

Transfer learning



Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.713	0.901	138,357,544	23
VGG19	549 MB	0.713	0.900	143,667,240	26
ResNet50	98 MB	0.749	0.921	25,636,712	-
ResNet101	171 MB	0.764	0.928	44,707,176	-
ResNet152	232 MB	0.766	0.931	60,419,944	-
ResNet50V2	98 MB	0.760	0.930	25,613,800	-
ResNet101V2	171 MB	0.772	0.938	44,675,560	-
ResNet152V2	232 MB	0.780	0.942	60,380,648	-
InceptionV3	92 MB	0.779	0.937	23,851,784	159
InceptionResNetV2	215 MB	0.803	0.953	55,873,736	572
MobileNet	16 MB	0.704	0.895	4,253,864	88
MobileNetV2	14 MB	0.713	0.901	3,538,984	88
DenseNet121	33 MB	0.750	0.923	8,062,504	121
DenseNet169	57 MB	0.762	0.932	14,307,880	169
DenseNet201	80 MB	0.773	0.936	20,242,984	201
NASNetMobile	23 MB	0.744	0.919	5,326,716	-
NASNetLarge	343 MB	0.825	0.960	88,949,818	-
EfficientNetB0	29 MB	-	-	5,330,571	-
EfficientNetB1	31 MB	-	-	7,856,239	-
EfficientNetB2	36 MB	-	-	9,177,569	-
EfficientNetB3	48 MB	-	-	12,220,535	-
EfficientNetB4	75 MB	-	-	19,466,823	-
EfficientNetB5	118 MB	-	-	30,562,527	-
EfficientNetB6	166 MB	-	-	43,265,143	-
EfficientNetB7	256 MB	-	-	66,658,687	-



Data Augmentation

Make the most of few training examples !
& Reduce Overfitting

K Keras

```
from keras.preprocessing.image import ImageDataGenerator  
  
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    rescale=1./255,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```

Image rotation



Image Zooming



Nearest
Reflect
Wrap
Constant

Image Shifting



Image Brightness



Data Augmentation

Make the most of few training examples !
& Reduce Overfitting

K Keras

```
from keras.preprocessing.image import ImageDataGenerator  
  
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    rescale=1./255,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```

Image Shear Intensity



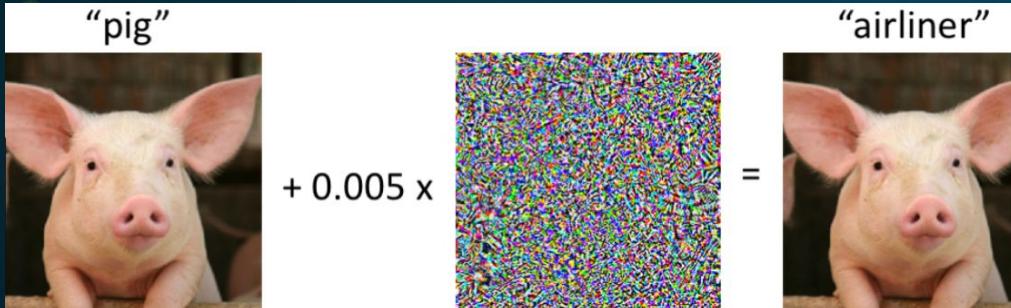
Horizontal flip



Robustness

03

Accuracy vs robustness



- The accuracy is not always the right metric
- Adding some well-chosen noise to a picture can completely

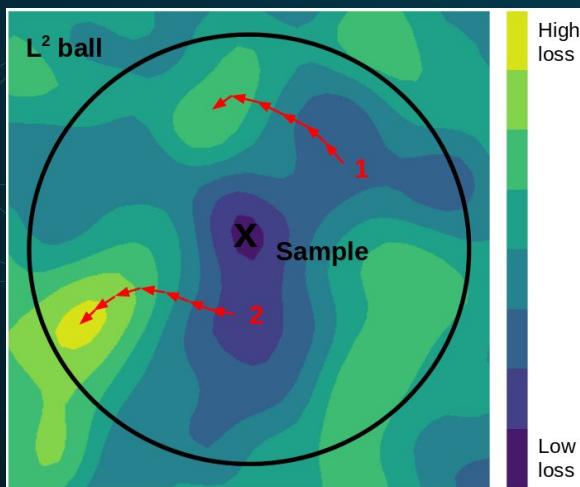
Fast Gradient Sign Method (FGSM)

- Introduced by Goodfellow and al. Explaining and Harnessing Adversarial Examples (ICLR 2015)
- Idea: compute the sign of the gradient of the loss wrt to each pixel of the input image

$$\begin{array}{ccc} \text{} & + .007 \times & \text{} \\ \mathbf{x} & & \text{sign}(\nabla_{\mathbf{x}} J(\theta, \mathbf{x}, y)) \\ \text{"panda"} & & \text{"nematode"} \\ 57.7\% \text{ confidence} & & 8.2\% \text{ confidence} \\ & = & \\ & & \text{} \\ & & \mathbf{x} + \\ & & \epsilon \text{sign}(\nabla_{\mathbf{x}} J(\theta, \mathbf{x}, y)) \\ & & \text{"gibbon"} \\ & & 99.3 \% \text{ confidence} \end{array}$$

Projected Gradient Descent (PGD)

- Introduced in Kurakin et al. 2016
- Idea: Apply FGSM several times while ensuring we stay in an ball around the original image



More complex attacks: Carlini & Wagner

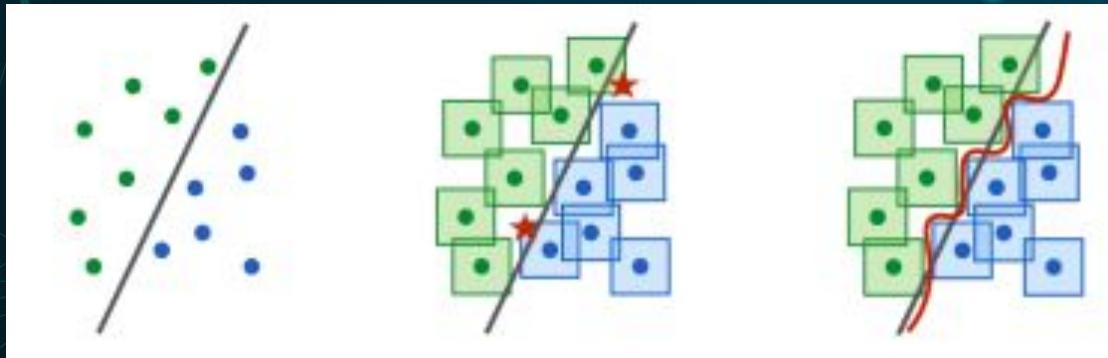
- Introduced by Carlini and Wagner, Towards Evaluating the Robustness of Neural Networks, IEEE 2017
- Idea: For a given example $x \in X$ of the class $y \in Y$, the l2 Carlini & Wagner attack (C&W) aims at solving the following optimization problem :

$$\min_{x+\delta} c\|\delta\|_2 + g(x + \delta)$$

- where $g(x + \delta) \leq 0$ if $f(x + \delta) = y$

The defense: How to make a model Robust

- A simple but yet effective way to defend against attacks is to add attacked images to the training set
- Introduced by Madry, Aleksander et al. , Towards deep learning models resistant to adversarial attacks (2017)



Generative Models

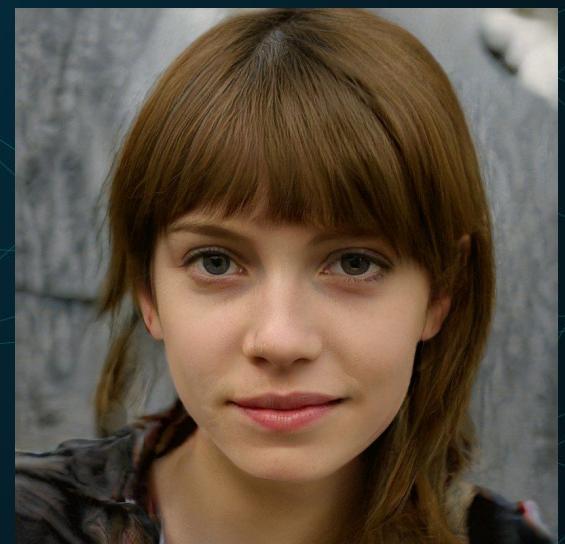
04

Gans

Generative models

- Models can detect, classify, but also generate.

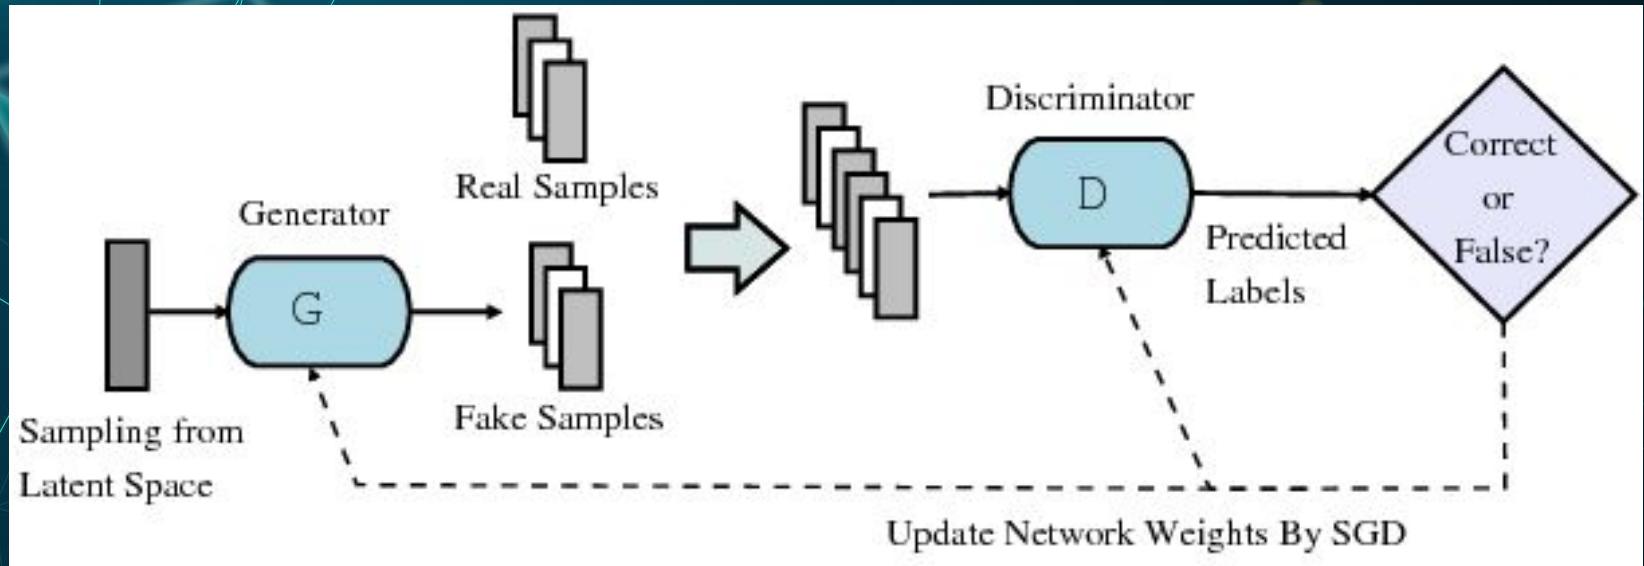
Image generated by StyleGAN, based on portrait analysis.



What is a GAN?

- Generative Adversarial Networks Introduced by Goodfellow and al. (2014)
- New approach: using Game Theory for training an image synthesis model
- Solving a minimax problem between a **Generator** and a **Discriminator**

Architecture of a GAN



source: <https://www.semanticscholar.org/paper/Generative-Adversarial-Networks-for-Unsupervised-Plakias-Boutalis/627d4f69b76bb3fc88283de9e8e9f7ee6c598ea7>

The optimisation problem - classical GANs

The Discriminator is trained to maximize:

$$L = E[\log P(S = \text{real} | X_{\text{real}})] + E[\log P(S = \text{fake} | X_{\text{fake}})]$$

⇒ Log-likelihood that the Discriminator assigns the generated image to the correct source

The Generator is trained to minimize:

$$E[\log P(S = \text{fake} | X_{\text{fake}})]$$

⇒ Second term in the first loss. Minimizing the log-likelihood that the Discriminator detects the fake images.

THANKS!

Do you have any questions?

olivier.randavel@gmail.com

hippolyte.mayard@dauphine.eu



CREDITS: This presentation template was created by **Slidesgo**, including icons by **Flaticon**, and infographics & images by **Freepik**.

Please keep this slide for attribution.