



Adversarial Attacks

Master Project

Team Members



Amin El Iraki



Olivier Randavel



Hadrien Mariaccia



Kenza Hammou

Github



The background features a teal-to-purple gradient with a pattern of white and light blue hexagons. Some hexagons are solid, while others are outlined. A network of thin white lines connects various points, some of which are marked with small teal dots. The overall aesthetic is modern and technical.

Evaluating several adversarial attacks

Part 2

Carlini L_∞



The optimization problem :

$$\begin{cases} \text{Minimize } D(x, x + \delta) \\ \text{st } C(x + \delta) = t, x + \delta \in [0, 1]^n \end{cases} \longrightarrow \text{which is not linear}$$

[f an objective function such that $C(x + \delta) = t$, if and only if $f(x + \delta) \leq 0$]

$$\begin{cases} \text{Minimize } D(x, x + \delta) \\ \text{st } f(x + \delta) \leq 0, x + \delta \in [0, 1]^n \end{cases}$$
$$\begin{cases} \text{Minimize } D(x, x + \delta) + c * f(x + \delta) \\ \text{st } x + \delta \in [0, 1]^n \end{cases}$$

$f_6(x') = (\max_{i \neq t} (Z(x')_i - Z(x')_t))^+$, Z softmax function

The L_∞ distance : $\|x - x'\|_\infty = \max(|x_1 - x'_1|, \dots, |x_n - x'_n|)$

The library used is **CarliniLInfMethod** with several parameters :

- epsilon : modification's pixel's limit
- learning_rate : SGD optimization coefficient
- max_iter : the number of iterations



Basic Iterative Method

=> an iterative version of FGSM




- Fast Method :
$$X^{adv} = X + \epsilon \text{sign}(\nabla_X J(X, y_{true}))$$
- Iterative Version :
$$X_0^{adv} = X$$
$$X_{N+1}^{adv} = \text{Clip}_{X, \epsilon} \left\{ X_N^{adv} + \alpha \text{sign}(\nabla_{X_N^{adv}} J(X, y_{true})) \right\}$$

The library used is **BasicIterativeMethod** with several parameters :

- epsilon : perturbation's limit
- eps_step : input variation at each iteration
- max_iter : the number of iterations

Newton Fool



The algorithm approximates the nearest classification boundary based on Newton's method for solving nonlinear equations.

1. The algorithm takes an input x_0 and a neural network with a softmax layer F_s such that $F_s(x_0) = l, \eta, \max_{l \neq \eta}$
2. Outputs the minimal perturbation required to misclassify the image

It tries to find small d_i to decrease the value of the function F as fast as possible to 0 : $F(x_0+d)_l \approx 0$.

Starting with x_0 , they approximated $F(x_i)$ using a linear function step by step as follows.

$$F_s^l(x) \approx F_s^l(x_i) + \nabla F_s^l(x_i)(x - x_i) \quad i \in [0, n] \rightarrow (1)$$

The library used is **NewtonFool** with several parameters :

- η : characterizes the perturbation
- \max_iter : the maximum number of iterations


```
!pip install adversarial-robustness-toolbox
from art.attacks import NewtonFool
classifier = KerasClassifier(model=model, clip_values=(0, 1))
attack_cw = NewtonFool(classifier=classifier, eta=0.1, max_iter=40,
targeted=False, batch_size=1)
x_test_adv = attack_cw.generate(x_test)
loss_test, accuracy_test = model.evaluate(x_test_adv, y_test)
perturbation = np.mean(np.abs((x_test_adv - x_test)))
```

Evaluation metrics

La perturbation : $\frac{1}{n} \sum_{i=1}^n |x_{i_{attacktest}} - x_{i_{test}}|$

L'évaluation des résultats : **model.evaluate(x_test_adv, y_test)**

L'évaluation des résultats : **show_dataset_and_predictions(x_test_adv, y_test, model, class_to_name=class_to_name)**



Results

	Output 6 pictures						x_test attack	avg perturbation	delta	epsilon	eps_step	max_iter	learning rate	batch size	targeted
FGSM							45.9%	0.010	0.008						
							39.7%	0.010	0.01						
							3.9%	0.080	0.1						
PGD L_infini							7.0%	0.030	0.01	0.008				2	
							4.4%	0.070	0.008	0.1				2	
							3.3%	0.070	0.008	0.008				5	
CarliniLinf							43.0%	0.001		0.03				4	0.01
							23.0%	0.001		0.03				40	0.01
							11.8%	0.010		0.25				100	0.3
Newton Fool							12.3%	0.004		0.01				100	
							12.4%	0.004		0.01				40	1
							13.1%	0.010		0.03				40	1
Basic Iterative Method							24.4%	0.010		0.03	0.01	2		1	FALSE
							10.1%	0.020		0.03	0.03	40		1	FALSE
							7.6%	0.070		0.1	0.03	40		1	FALSE



PGD L_∞


FGSM


Basic
Iterative
Method

CarliniL_∞

Newton
Fool

Perturbations distribution per attack





References

Carlini L_{∞} : Towards Evaluating the Robustness of Neural Networks - Nicholas Carlini & David Wagner

<https://arxiv.org/pdf/1608.04644.pdf>

Newton Fool : Objective Metrics and Gradient Descent Algorithms for Adversarial Examples in Machine Learning

<https://andrewxiwu.github.io/public/papers/2017/JWJ17-objective-metrics-and-gradient-descent-based-algorithms-for-adversarial-examples-in-machine-learning.pdf>

Basic Iterative Method : Adversarial Machine Learning At Scale,

<https://arxiv.org/pdf/1611.01236.pdf>

Adversarial Examples in the Physical World

<https://arxiv.org/pdf/1607.02533.pdf>

Latex maths : <http://latex2png.com/>



THANKS